

## Message flow in kafka:



### 0. Core Concepts Before We Begin

- Producer: Publishes messages to Kafka.
  - Consumer: Reads messages from Kafka.
  - Topic: Logical grouping/category to which messages are sent.
  - Broker: Kafka server that stores data and serves clients.
  - Partition: Each topic is split into one or more partitions for parallelism and scalability.
  - Offset: A unique identifier for each message within a partition.
  - Consumer Group: A group of consumers cooperating to read from a topic.
- 



### 1. Message Creation & Publishing (Producer Side)

#### Step-by-step:

1. **Message Created: A producer creates a message, usually consisting of:**
  - Key (optional): Used for partitioning logic.
  - Value: The actual payload (event data).
  - Headers / Metadata (optional).
2. **Topic Decided: The producer chooses a topic (e.g., "orders") to send the message to.**
3. **Partition Decision:**
  - If a key is provided, Kafka uses a hashing algorithm (e.g., `hash(key) % number_of_partitions`) to decide which partition.
  - If no key, Kafka may use round-robin or random distribution.

#### 4. Message Sent to Broker:

- Kafka client sends the message to the leader broker for that partition.
  - Kafka writes the message sequentially to a commit log file (append-only) on disk.
- 

## 2. Message Storage (Kafka Broker Side)

- Each partition is a log file on disk.
  - Message is written with a monotonic offset, unique within that partition (starting from 0).
  - Kafka does not delete messages after consumption. Instead, it retains messages for a configurable period (e.g., 7 days).
  - Brokers replicate partitions across other brokers for fault tolerance (controlled by replication factor).
- 

## 3. Message Consumption (Consumer Side)

### Consumer Mechanics:

1. Consumer subscribes to a topic (e.g., "orders").
  2. Kafka assigns specific partitions to each consumer (within a consumer group).
  3. Consumers read messages sequentially from the assigned partition using the offset.
- 

## 4. Offset Management

- Offset is the position of a message in a partition.
- Consumers maintain an offset per partition to track how far they have read.
- Two types of offset storage:

- Automatic (Kafka managed): Kafka stores it in a special topic (`__consumer_offsets`).
  - Manual (App managed): Developers store offsets in external systems or commit them manually.
  - **Offsets can be committed:**
    - Automatically after consumption.
    - Manually, based on business logic (e.g., after processing is successful).
- 

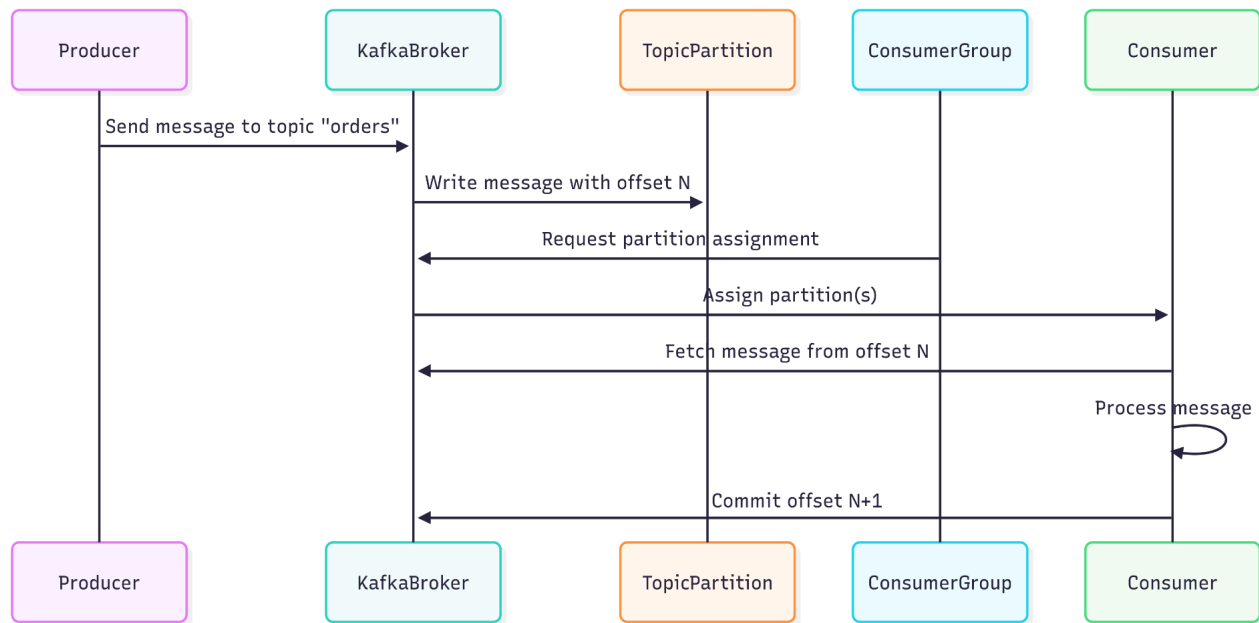
## 5. Multiple Consumers (Consumer Groups)

### Key Concepts:

- A consumer group allows multiple consumers to read from the same topic in parallel.
- Kafka guarantees each partition is read by only one consumer within a group.

### Scenarios:

- If there are more partitions than consumers, some consumers get multiple partitions.
- If there are more consumers than partitions, some consumers remain idle.
- If each consumer is in a separate group, each one gets the full copy of all topic data (like pub-sub broadcast).



---

## 🔧 6. Partition Rebalancing

- When a consumer joins or leaves a group, Kafka rebalances partitions.
- Rebalancing involves pausing consumption temporarily and redistributing partitions.

---

## 🧠 7. Fault Tolerance and Durability

- Kafka's replication factor ensures data isn't lost if a broker crashes.
- Consumers can resume from last committed offset after crash/restart.
- Producers get ACKs based on configured `acks=0/1/all` for durability guarantees.

## 🔍 Intricacies You Should Know

Intricacy	Detail
<b>Message Ordering</b>	Guaranteed only <b>within a partition</b> , not across partitions.
<b>Backpressure Handling</b>	If consumer is slow, offset lag increases, but Kafka doesn't block producer.
<b>Compaction</b>	Kafka can <b>compact</b> messages with the same key (latest retained).
<b>Retention Policies</b>	Based on time (e.g., 7 days) or size (e.g., 1GB per partition).
<b>High Throughput</b>	Due to zero-copy, batch writes, and sequential disk I/O.
<b>Exactly Once Semantics (EOS)</b>	Supported with proper configuration: idempotent producers, transactions, and offset coordination.

---

### ✓ TL;DR (In One Line Each Stage)

1. Producer creates a message and sends it to a topic.
2. Kafka decides partition via key hashing or round-robin.
3. Broker writes the message to the partition log with an offset.
4. Consumer reads from the topic by partition and offset.
5. Offsets track progress; committed offsets let consumers resume.

6. Consumer groups ensure parallelism with one consumer per partition.
7. Kafka manages rebalancing, fault tolerance, and retention.

## Apache Kafka vs RabbitMQ

---

### Apache Kafka

Kafka is an open-source distributed event streaming platform developed by the Apache Software Foundation. It is written in Java and Scala. Kafka uses a data stream for the delivery of messages and is suitable for both online and offline message consumption. The aim of the project is to provide a high-throughput, low-latency platform for real-time data feed handling, streaming analytics, data integration, and mission-critical applications. It employs a pull-based approach and lets users request message batches from specific offsets. Although only being shipped with a Java client, it offers an adapter SDK that allows programmers to build their own system integrations.

#### Features of Kafka

- High throughput: Built to process millions of messages per second and handle large volumes of data.
  - Zero downtime and data loss: Achievable with a replication factor  $> 1$  for brokers.
  - Ingest pipelines: Kafka can replicate events in a broker.
  - Reliable: Due to its distributed, replicated, partitioned, and fault-tolerant nature.
- 

### RabbitMQ

RabbitMQ is an open-source distributed general-purpose message-broker software facilitating efficient message delivery in complex routing scenarios. It runs as a cluster of nodes where queues are distributed across them.

It initially implemented the **Advanced Message Queuing Protocol (AMQP)**. It uses a **push model** and prevents overwhelming consumers via the consumer-configured prefetch limit. It is commonly used for handling background jobs or as a message broker between microservices.

## Features of RabbitMQ

- Built-in clustering: Multiple servers can be clustered into a single broker.
- Protocol support: Supports AMQP, STOMP, MQTT, and more.
- Tracing and debugging: Allows investigation when the system misbehaves.
- Broad language support: Clients available for most programming languages.

---

## Note

Kafka and RabbitMQ serve the same core purpose—event/message handling—but with different strengths. Both are open-source, commercially supported pub/sub systems widely used by enterprises.

---

## Kafka vs RabbitMQ: Comparison Table

S.No	Parameter	Kafka	RabbitMQ
1	Performance	Around 1 million messages per second	Around 4k–10k messages per second
2	Payload Size	Default limit: 1 MB	No strict size constraints
3	Messages Synchronicity	Durable store, allows replay	Can be asynchronous/synchronous
4	Data Unit	Continuous stream	Message
5	Data Type	Operational	Transactional
6	Topology	Publish/Subscribe based	Exchange types: Fanout, Topic, Direct, Header-based
7	Data Flow	Continuous unbounded key-value pairs	Bounded distinct message packets

8	Event Storage Structure	Logs	Queues
9	Message Delivery System	Pull-based (consumer requests messages)	Push-based (messages pushed to queues)
10	Consumer Mode	Smart consumer / Dumb broker	Dumb consumer / Smart broker

---

## When to Use Kafka

- When very fast message consumption is required.
- When message replay is needed.
- For high throughput requirements (100K+ events/sec).
- Suitable for real-time analytics, event sourcing, log aggregation, etc.

## When to Use RabbitMQ

- When dynamic consumer addition is needed (without publisher changes).
- If message replay is not necessary.
- When supporting legacy protocols like AMQP 0-9-1, STOMP, AMQP 1.0, and MQTT.
- Ideal for legacy systems and complex routing logic.